

Ryan Choi

Senior Full Stack AI Engineer

✉ leeson.nerd@gmail.com ☎ +1 312 680 2246 📍 Brooklyn, NY

🌐 <https://linkedin.com/in/ryan-soon-nerd>

PROFESSIONAL EXPERIENCE

Software Engineer, Intellectyx, Inc

04/2025 – Present | Brooklyn, NY

Project: LLM-powered Knowledge Assistant (0→1 build for enterprise document intelligence)

- Took a **0→1 AI product from idea to production**, building an LLM-powered knowledge assistant for semantic search, document Q&A, and workflow automation.
- Designed and shipped a **RAG pipeline end-to-end** (chunking → embeddings → retrieval → generation) using OpenAI/Hugging Face + Pinecone/FAISS, significantly improving answer relevance over naive prompting.
- Wrote scrappy but scalable **data pipelines for messy, unstructured enterprise data** (PDFs, CSVs, internal docs), making them queryable via natural language.
- Built the entire backend in FastAPI, handling ingestion, embedding pipelines, and low-latency inference APIs under real-world usage.
- Implemented **streaming LLM responses + async workflows**, making the product feel fast and interactive despite heavy model calls.
- Built a lightweight **React/Next.js frontend** to demo and ship features fast, including chat UI and document-level insights.

Senior Full Stack AI Engineer, Suki AI

03/2022 – 04/2025 | Brooklyn, NY

Focus: AI-powered clinical assistant (LLMs + speech + healthcare workflows)

- Led development of **AI-driven clinical assistant features**, enabling doctors to automate documentation, extract insights, and interact with patient data using natural language.
- Built and scaled **LLM-powered pipelines** for clinical note generation, summarization, and entity extraction, improving accuracy and usability in real-world healthcare settings.
- Designed and optimized **RAG-based systems** integrating EHR data, medical knowledge bases, and vector search to deliver context-aware responses.
- Worked extensively on **speech-to-text + NLP integration**, transforming raw audio into structured, actionable clinical documentation.
- Improved model performance via **prompt engineering, fine-tuning strategies, and evaluation pipelines**, ensuring outputs met healthcare quality standards.
- Ensured **HIPAA-compliant data handling**, incorporating strict security, auditing, and access control mechanisms.

Software Engineer, Valvoline

02/2020 – 02/2022 | Brooklyn, NY

Project: Vehicle Product Recommendation & Service Platform (used for product selection + maintenance workflows)

- Built and shipped features for a **vehicle-based product recommendation platform**, enabling users to find the right motor oil, fluids, and maintenance products based on vehicle data.
- Developed backend services (Java/Python) powering **product lookup, compatibility logic, and recommendation workflows** across a large catalog of automotive products.
- Designed and optimized **data models and APIs** to support product categories like motor oils, coolants, and transmission fluids, ensuring fast and accurate results.
- Automated data ingestion and validation pipelines for **vehicle specifications and product compatibility datasets**.
- Debugged production issues and improved system reliability for applications supporting **nationwide service operations**.
- Collaborated with product and business teams to translate requirements into **scalable, customer-facing digital features**.

EDUCATION

Bachelor's Degree in Computer Science, Cornell University

2016 – 2020 | Ithaca, NY